



703308 VO High-Performance Computing MPI Derived Datatypes and Virtual Topologies

Philipp Gschwandtner

Overview

- ▶ derived datatypes
 - ▶ allows to send user-specific datatypes
- ▶ virtual topologies
 - ▶ adds semantic position information to ranks
- ▶ tales from the proseminar

Motivation

- ▶ we discussed using MPI for parallelization, but on a very basic level
 - ▶ we can only transfer contiguous ranges of arrays of the same element type
 - ▶ we need to manually compute rank numbers for talking to semantically significant and often-used ranks (e.g. left/right neighbor)
- ▶ what about
 - ▶ transferring (nested) structs/classes, arrays of tuples, columns of a 2 D matrix, etc.
 - ▶ ease of coding & semantics: “send temperature to my left neighbor” instead of “send double to $(\text{myRank} - 1 + \text{numRanks}) \% \text{numRanks}$ ”



Derived Datatypes



Recap: MPI datatypes

- ▶ several predefined basic types
 - ▶ MPI_INT, MPI_FLOAT, MPI_DOUBLE, MPI_BYTE, ...
- ▶ what about something like on the right?
 - ▶ struct with 4 members
 - ▶ 3x 8 bytes + 4 bytes

```
struct Particle {  
    double x;  
    double y;  
    double z;  
    int species;  
};
```

Issues with more complex data structures

- ▶ MPI doesn't know how large a single element is
 - ▶ no predefined `MPI_(DATA_TYPE_THAT_DOESNT_EXIST_YET)`
 - ▶ what about nesting types? with differently-sized members?
 - ▶ sending individual elements blows up the code and causes performance overhead due to multiple messages
- ▶ issue of sending a single member of struct instances
 - ▶ bad solution: explicitly assemble send/receive buffers with single data type per message transfer
 - ▶ causes coding, memory footprint, and message overhead (at least one message per type)

Why not just use MPI_BYTE everywhere?

- ▶ derived datatypes add type information, allow automatic type handling
 - ▶ size of e.g. `int` is unknown (C standard only defines minimum requirements!)
 - ▶ `int` on machine A and `int` on machine B might have different size
 - ▶ machine A might be little-endian, machine B might be big-endian
 - ▶ saves a lot of explicit user-written `sizeof()` constructs
 - ▶ enables type-specific hardware optimizations for MPI
- ▶ using `MPI_BYTE/...` everywhere deprives you of all of the above
 - ▶ also does not carry any semantic information on the content
 - ▶ does not work for C++ objects that are not PODs (Plain Old Data)
 - ▶ undefined behavior due to pointers & vtables (+ more reasons depending on setup)

MPI derived datatypes

- ▶ composed of existing types
 - ▶ both basic and derived
 - ▶ can be nested
- ▶ used to transfer high-level data structures
 - ▶ encodes more information in transfer, allows MPI to perform optimizations
 - ▶ more performance-efficient than individual transfer of data structure contents
 - ▶ less code, easier to read and maintain

MPI derived datatypes cont'd

- ▶ allow definition of new handles
 - ▶ e.g. MPI_FOOBAR
- ▶ require several steps
 - ▶ construction: declare and define new datatype
 - ▶ allocation / commit: needs to be done once by all ranks before using new datatype
 - ▶ usage
 - ▶ deallocation: frees internal MPI storage, to be done once by all ranks

Selection of MPI derived datatype facilities

- ▶ `MPI_Type_create_struct(...)`
 - ▶ specifies the data layout of user-defined structs (or classes)
- ▶ `MPI_Type_vector(...)`
 - ▶ specifies strided data, i.e. same-type data with missing elements
- ▶ `MPI_Type_create_subarray(...)`
 - ▶ specifies sub-ranges of multi-dimensional arrays
- ▶ `MPI_Type_contiguous(...)`
 - ▶ specifies a user-defined contiguous type comparable to C arrays

Structs

- ▶ `int MPI_Type_create_struct(int count,
 const int blocklengths[], const MPI_Aint
displacements[],
 const MPI_Datatype types[], MPI_Datatype* newtype)`
 - ▶ `count`: number of blocks
 - ▶ `blocklengths`: number of elements per block (array)
 - ▶ `displacements`: starting address of first element of each block (array)
 - ▶ `types`: type of each block (array)
 - ▶ `newtype`: resulting derived datatype

Structs: block lengths, displacements and types

```
struct Particle {  
    int posX;  
    int posY;  
    int posZ;  
    double magneticForceX;  
    double magneticForceY;  
    double magneticForceZ;  
};
```

} block no 0, starts at byte 0,
3 elements of type integer

} block no 1, starts at byte 12,
3 elements of type double

Careful with displacements

- ▶ careful with manually specifying displacements
 - ▶ binary layout of structs in memory is compiler-dependent (e.g. struct members might be padded)
 - ▶ use `offsetof()` instead!
 - ▶ careful with C vs. C++ differences, e.g. size of empty structs/classes
- ▶ also, do not confuse `MPI_Aint` (programming language data type) and `MPI_AINT` (MPI data type)
- ▶ additional option: use `MPI_BOTTOM` as the buffer argument, enables use of absolute addresses as displacements instead of offsets
 - ▶ rationale: `MPI_Aint` is a signed integer (overflow behavior is not defined), absolute addresses are unsigned (overflow behavior is defined)

```
MPI_Aint displacements[2] =  
    { 0,  
      12 };  
  
// ===== vs =====  
  
MPI_Aint displacements[2] =  
    { offsetof(Foo, posX),  
      offsetof(Foo, magneticForceX) };
```

Careful with pointers

- ▶ don't transfer shallow copies of data
 - ▶ `double*` data might not be available or likely at a different memory address on node B
- ▶ try to avoid
 - ▶ otherwise, make a deep copy and adjust pointers at receiver side

```
struct Particle {  
    int size;  
    double* data;  
};
```

Struct example

```
typedef struct {
    int barInt;
    double barDoubleA;
    double barDoubleB;
} Foo;
MPI_Datatype myType;
int blocklengths[2] = { 1, 2 };
MPI_Aint displacements[2] =
    { offsetof(Foo, barInt),
      offsetof(Foo, barDoubleA) };
MPI_Datatype datatypes[2] =
    { MPI_INT, MPI_DOUBLE };
MPI_Type_create_struct(2, blocklengths,
    displacements, datatypes, &myType);
```

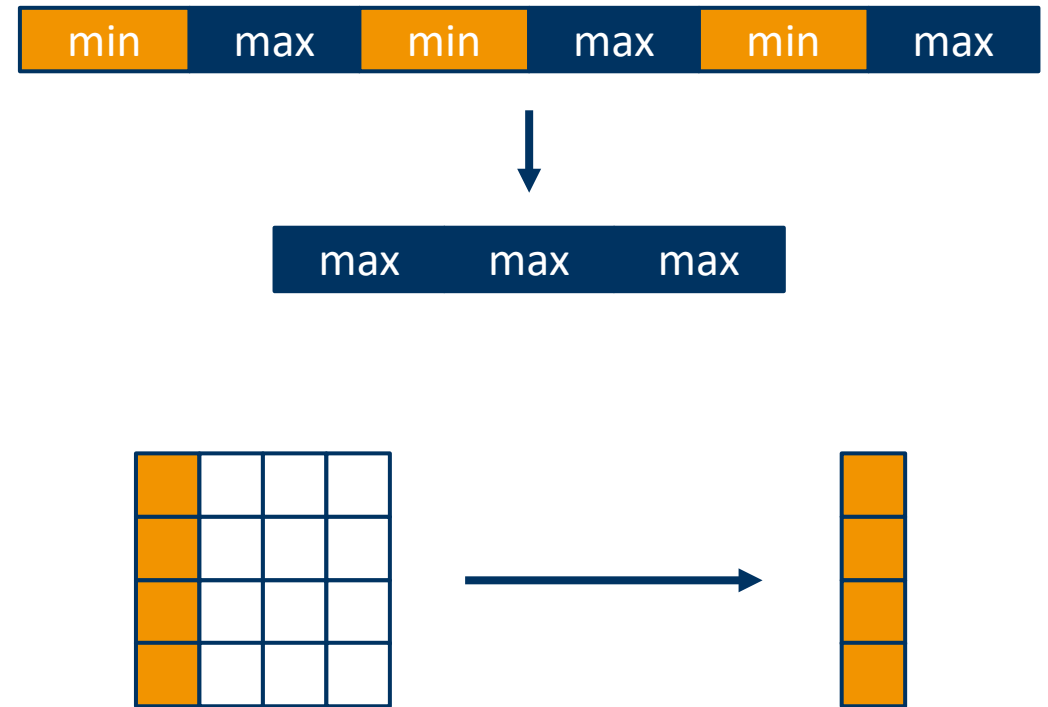
```
MPI_Type_commit(&myType);

if (myRank == 0) {
    Foo data[2] = ...
    MPI_Send(data, 2, myType, 1, 42,
        MPI_COMM_WORLD);
} else {
    Foo data[2] = ...
    MPI_Recv(data, 2, myType, 0, 42,
        MPI_COMM_WORLD,
        MPI_STATUS_IGNORE);
}

MPI_Type_free(&myType);
```

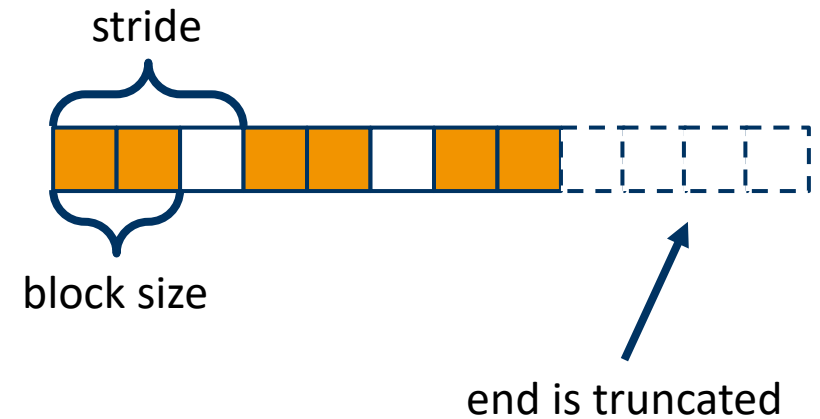
Non-contiguous data

- ▶ send all max values of an array of (min, max)-tuples to another rank
- ▶ send the column of a matrix
- ▶ do all of that without having to copy data to a contiguous buffer first!



Vectors

- ▶ Support strides (gaps in arrays)
 - ▶ e.g. take 2 elements, omit 1 element, repeat 3 times in total
 - ▶ useful for linear algebra



Vector example

```
#define SIZE 20
#define COUNT 3
#define LENGTH 2
#define STRIDE 3

MPI_Datatype myType;
MPI_Type_vector(COUNT, LENGTH, STRIDE,
               MPI_CHAR, &myType);
MPI_Type_commit(&myType);
```


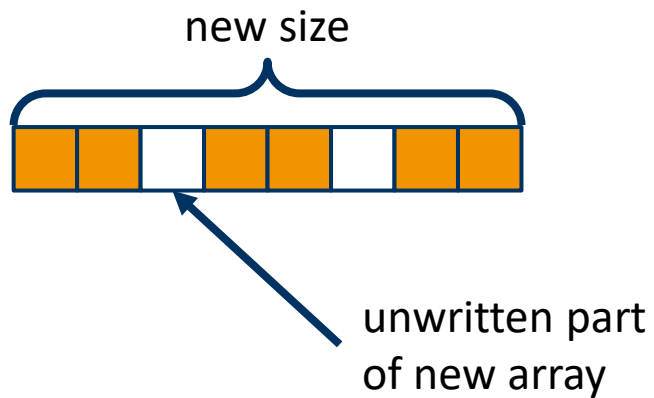
```
if (myRank == 0) {
    char data[SIZE] = ...;
    MPI_Send(data, 1, myType, 1, 42,
             MPI_COMM_WORLD);
} else {
    char data[SIZE];
    MPI_Recv(data, 1, myType, 0, 42,
             MPI_COMM_WORLD,
             MPI_STATUS_IGNORE);
}

MPI_Type_free(&myType);
```

Vector variants: Use case compaction

```
char data[SIZE];  
MPI_Recv(data, 1, myType,  
         0, 42, MPI_COMM_WORLD,  
         MPI_STATUS_IGNORE);
```

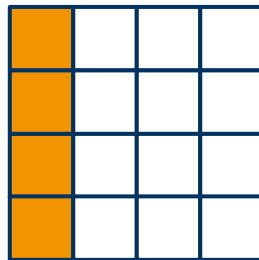
```
char data[COUNT*LENGTH];  
MPI_Recv(data, COUNT*LENGTH, MPI_CHAR,  
         0, 42, MPI_COMM_WORLD,  
         MPI_STATUS_IGNORE);
```



Vector variants: use case transposition

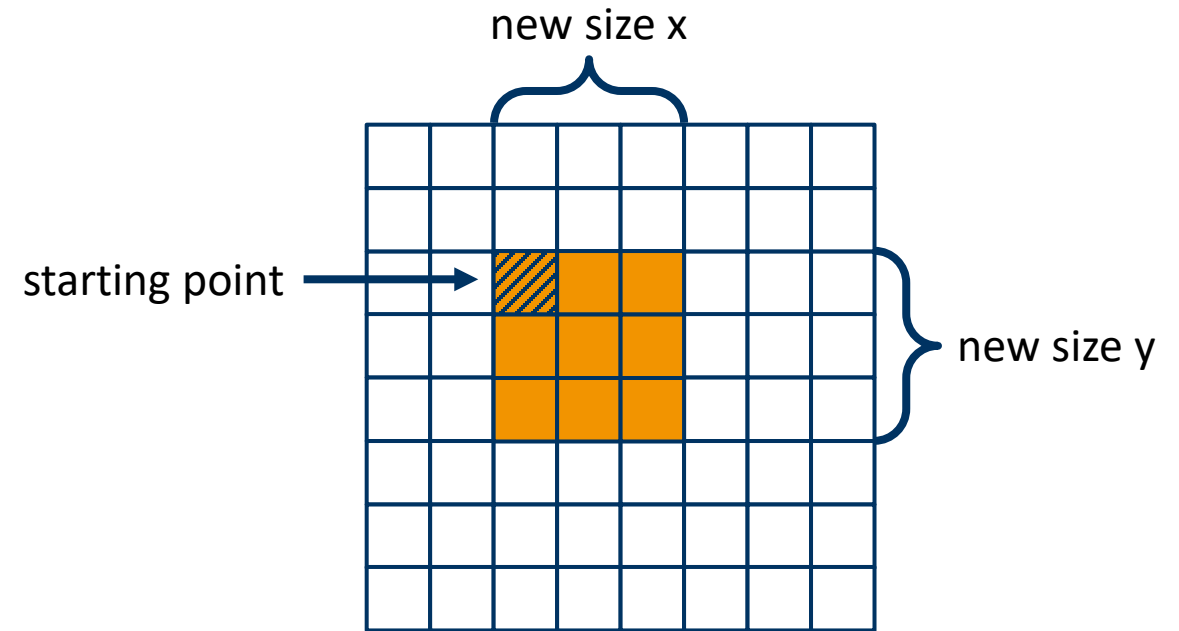
```
int data[SIZE][SIZE];
MPI_Type_vector(SIZE, 1, SIZE, MPI_INT,
                &myType);
MPI_Type_commit(&myType);
MPI_Send(data, 1, myType,
         1, 42,
         MPI_COMM_WORLD);
```

```
int data[SIZE];
MPI_Recv(data, SIZE, MPI_INT,
         0, 42,
         MPI_COMM_WORLD, MPI_STATUS_IGNORE);
```



Subarrays

- ▶ Allows to address a contiguous multi-dimensional sub-range of array elements



Subarray example in 2 D

```
#define SIZE 8
#define SUBSIZE 3

MPI_Datatype myType;
int size[2] = { SIZE, SIZE };
int subSize[2] = { SUBSIZE, SUBSIZE };
int start[2] = { 2, 2 };

MPI_Type_create_subarray(2, size,
    subSize, start, MPI_ORDER_C, MPI_INT,
    &myType);
MPI_Type_commit(&myType);
```

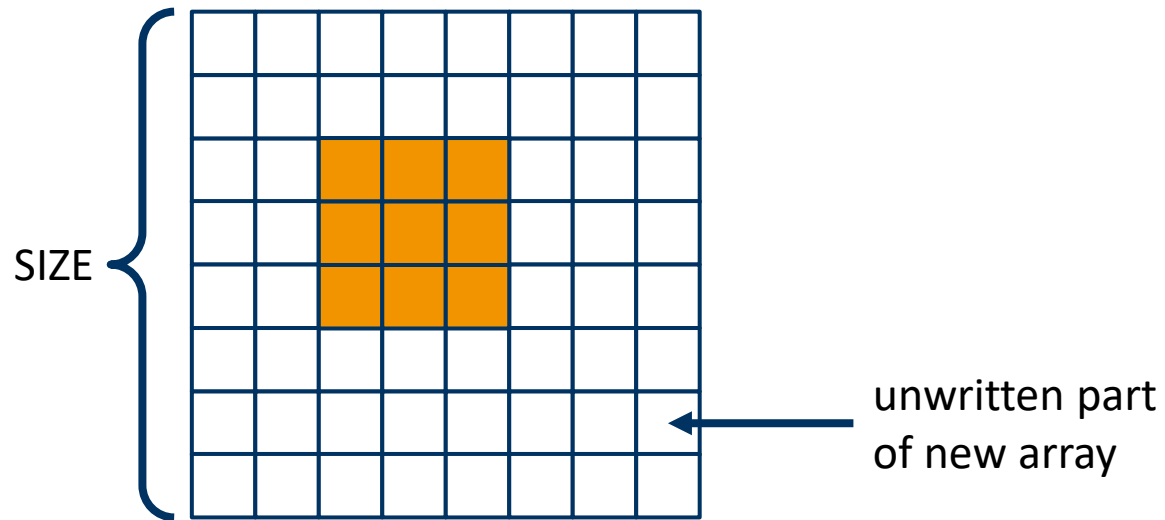
```
if (myRank == 0) {
    int data[SIZE][SIZE] = ...;
    MPI_Send(data, 1, myType, 1, 42,
        MPI_COMM_WORLD);
} else {
    int subData[SUBSIZE][SUBSIZE];
    MPI_Recv(subData, SUBSIZE*SUBSIZE,
        MPI_INT, 0, 42, MPI_COMM_WORLD,
        MPI_STATUS_IGNORE);
}

MPI_Type_free(&myType);
```

Subarray receive variants

```
int data[SIZE][SIZE];  
MPI_Recv(data, 1,  
         myType, 0, 42, MPI_COMM_WORLD,  
         MPI_STATUS_IGNORE);
```

```
int subData[SUBSIZE][SUBSIZE];  
MPI_Recv(subData, SUBSIZE*SUBSIZE,  
         MPI_INT, 0, 42, MPI_COMM_WORLD,  
         MPI_STATUS_IGNORE);
```



Multiple ways of distributing columns

- ▶ Allocate as a 1D array, use linearized indices
 - ▶ use 1D MPI vector with stride
 - ▶ (use nD MPI subarray with 1 dimension)
 - ▶ (use nD MPI darray with 1 dimension)
- ▶ Allocate as an nD array
 - ▶ use nested 1D MPI vectors
 - ▶ use nD MPI subarray
 - ▶ use nD MPI darray
- ▶ Same functional result for all of the above, but performance might differ
 - ▶ remember, MPI doesn't guarantee performance portability

Contiguous derived datatypes

- ▶ allows to aggregate same-type arrays into a single-count datatype
- ▶ has certain advantages
 - ▶ sending more than INT_MAX elements (count parameter type in MPI_Send/Recv/... is only int!)
 - ▶ allows semantic grouping and naming of data

```
MPI_Datatype myType;
MPI_Type_contiguous(SIZE, MPI_CHAR, &myType);
MPI_Type_commit(&myType);

char data[SIZE] = { 0 };

if(myRank == 0) {
    MPI_Send(data, 1, myType, 1, 42,
             MPI_COMM_WORLD);
} else {
    MPI_Recv(data, 1, myType, 0, 42,
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
}

MPI_Type_free(&myType);
```

Packing/unpacking

- ▶ MPI also offers `MPI_Pack(...)` and `MPI_Unpack(...)` functions
 - ▶ “Packs a datatype into contiguous memory” (MPICH documentation)
 - ▶ prefer this over derived datatypes? (hint: no)
- ▶ requires explicit copy of data from non-contiguous, user-defined form into a contiguous buffer to be sent with MPI
 - ▶ mostly superseded by MPI functions presented thus far, which may directly access user-defined structures (no user copy required)
 - ▶ pack/unpack still mostly offered for compatibility reasons
 - ▶ only very few modern edge cases

Free the datatypes!

- ▶ call `MPI_Type_free(...)` once you no longer need the type
 - ▶ frees MPI-internal data storage for your custom type
 - ▶ reduces memory footprint for large numbers of datatypes
 - ▶ facilitates debugging
 - ▶ note: any pending communication using this type will continue and complete normally
 - ▶ omitted in some source code examples on my slides for obvious space reasons

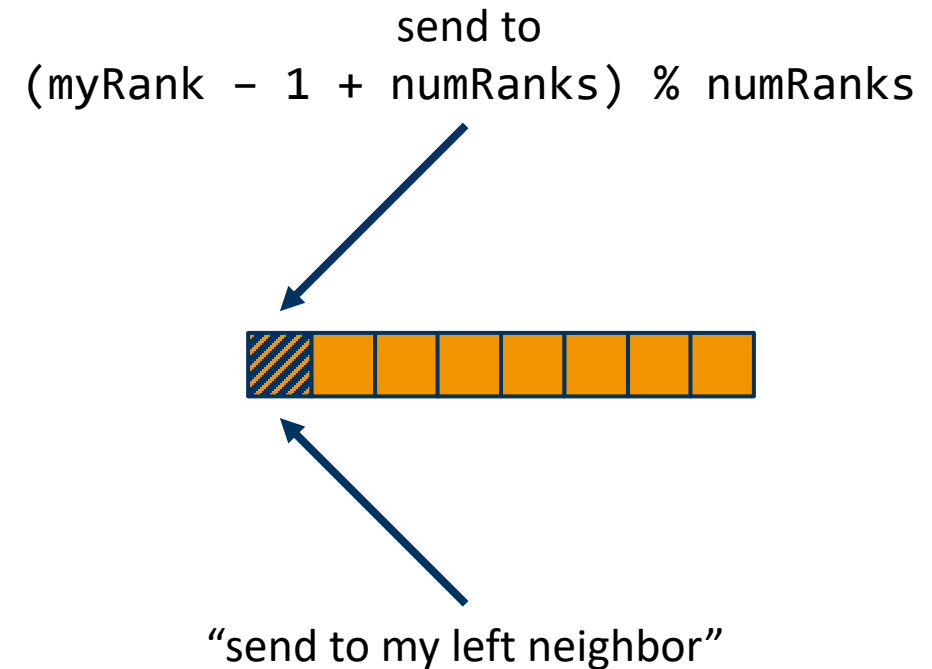


Virtual Topologies



Virtual topologies


- ▶ allows to “name” MPI ranks and provide addresses with semantics
 - ▶ high-level view of MPI ranks
 - ▶ simplifies implementation of complex algorithms
 - ▶ called “virtual” because it’s independent of the hardware topology
- ▶ naming scheme should fit communication pattern
 - ▶ and reflect the real-world topological relationship of parts of your problem
 - ▶ enables MPI to perform optimizations



There are two types of topologies (according to MPI)


▶ Cartesian topologies

- ▶ regular grids of squares/cubes/...
- ▶ each rank is a node on the grid and connected to its neighbors
- ▶ boundaries can be periodic
- ▶ ranks can be identified via Cartesian coordinates instead of rank ID

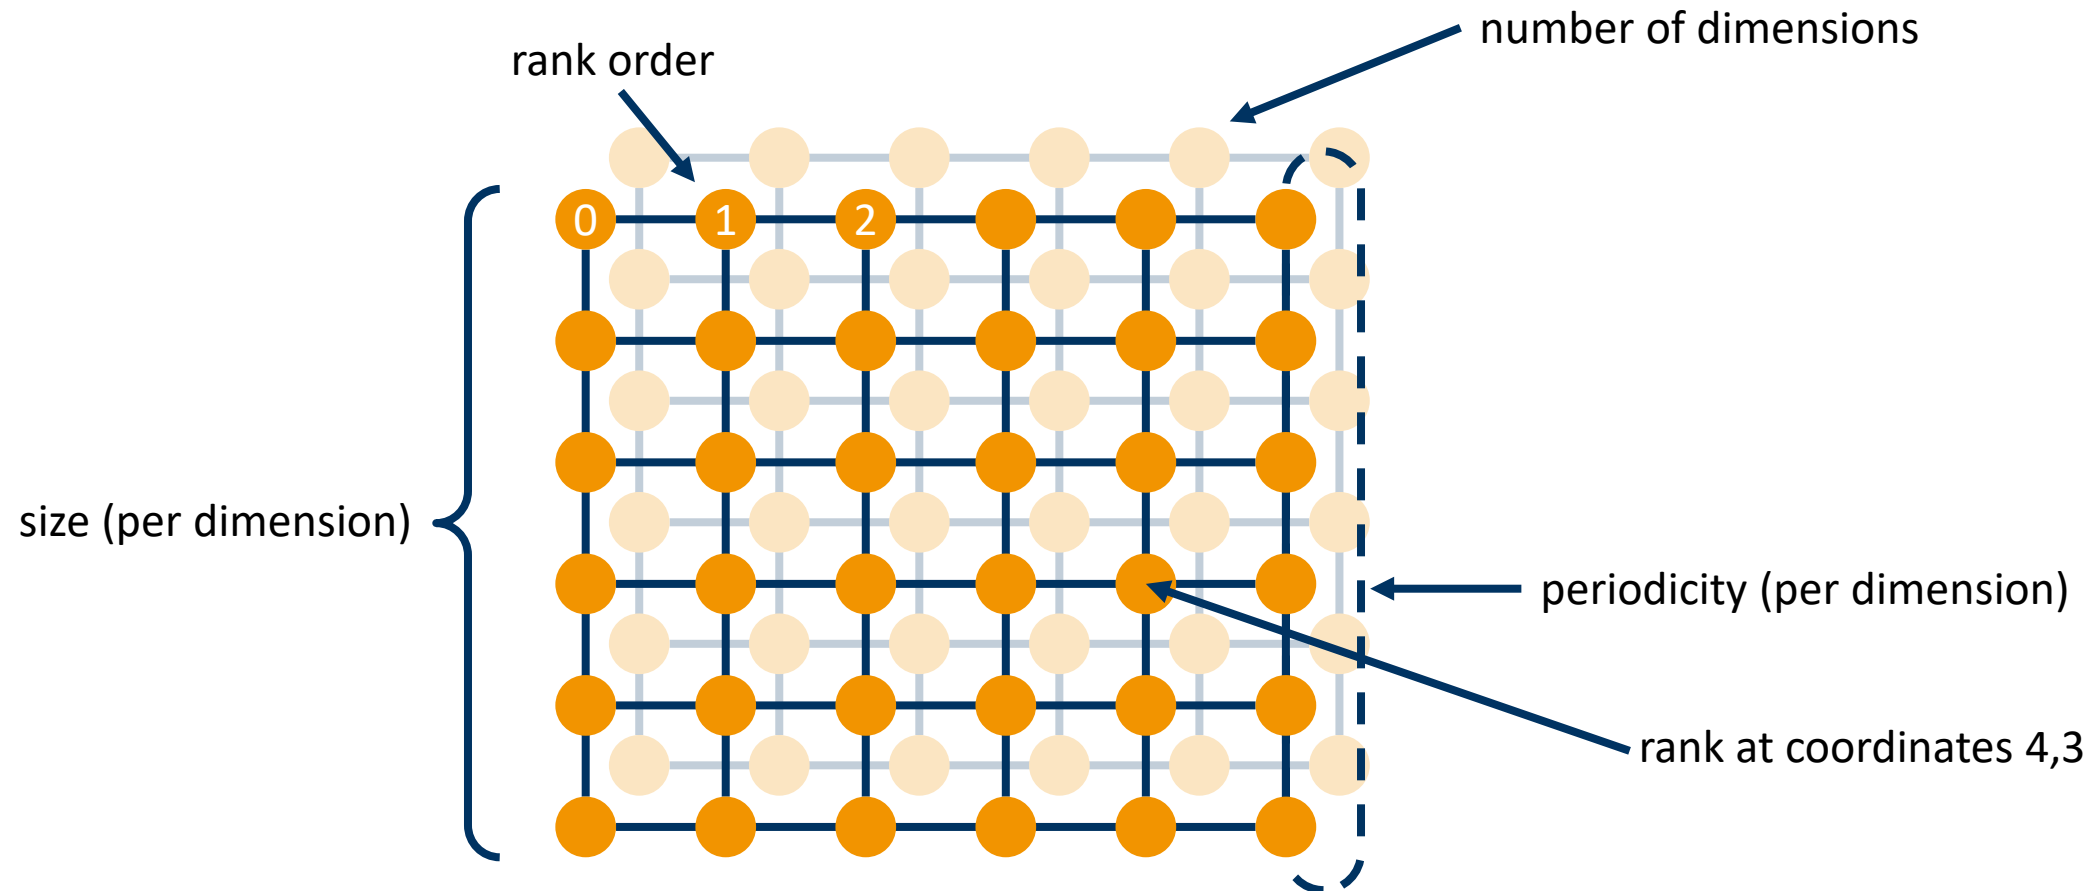


▶ graph topologies

- ▶ general graphs
- ▶ each rank is a vertex in the graph
- ▶ edges represent neighbor relationship
- ▶ edge weights specify communication intensity (facilitates optimization)
- ▶ not covered here



Properties of Cartesian topologies



Working with Cartesian topologies

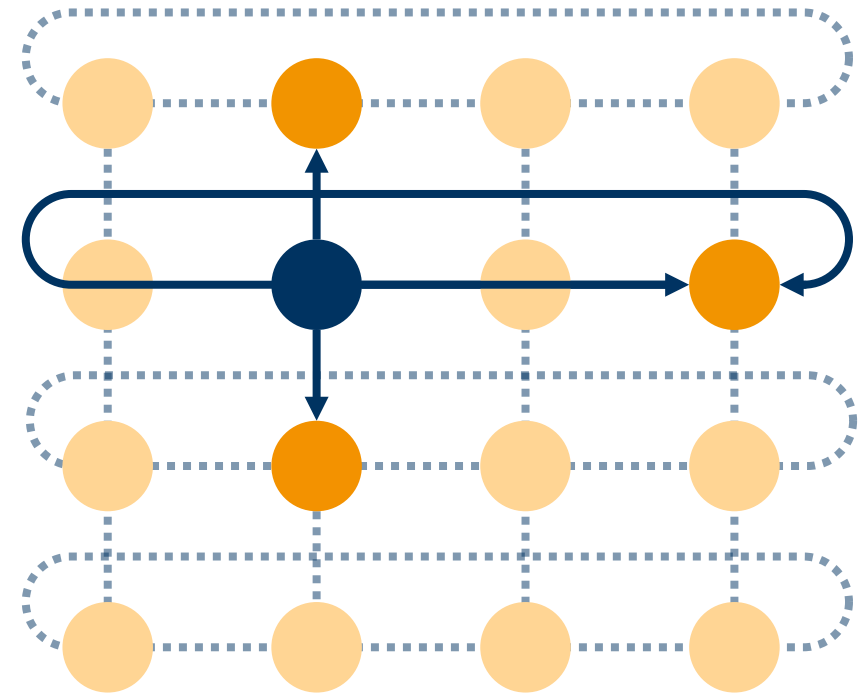
- ▶ create topology, resulting in new communicator
 - ▶ need to decide on dimensions, sizes, periodicity, etc...
 - ▶ per-dimension sizes can be computed using convenience function `MPI_Dims_create()`
 - ▶ new communicator implies rank IDs might have changed!
 - ▶ (remember MPI basics lecture: “[...] MPI semantics are relative to a “*communicator*” or “*group*”)
- ▶ (re)compute rank numbers or coordinates as required
- ▶ communicate as you please
 - ▶ remember to specify correct communicator from this point on

Creating a Cartesian topology

- ▶ `int MPI_Cart_create(MPI_Comm comm_old, int ndims, const int dims[], const int periods[], int reorder, MPI_Comm* comm_cart)`
 - ▶ `comm_old`: current communicator
 - ▶ `ndims`: number of dimensions
 - ▶ `dims`: size, per dimension
 - ▶ `periods`: periodicity (0 = open, 1 = periodic), per dimension
 - ▶ `reorder`: reorder rank numbers (0 = false, 1 = true)
 - ▶ `comm_cart`: new communicator with cartesian topology


Shifting (`MPI_Cart_shift()`)

- ▶ computes rank numbers of neighbors
 - ▶ requires direction and displacement (=distance)
- ▶ example on the right
 - ▶ partially periodic 2D topology of 4x4
 - ▶ up/down shift with displacement 1
 - ▶ or left/right shift with displacement 2
- ▶ Can return `MPI_PROC_NULL` if neighbor does not exist
 - ▶ can be used in communication, will result in a no-op



Slicing (`MPI_Cart_sub()`)

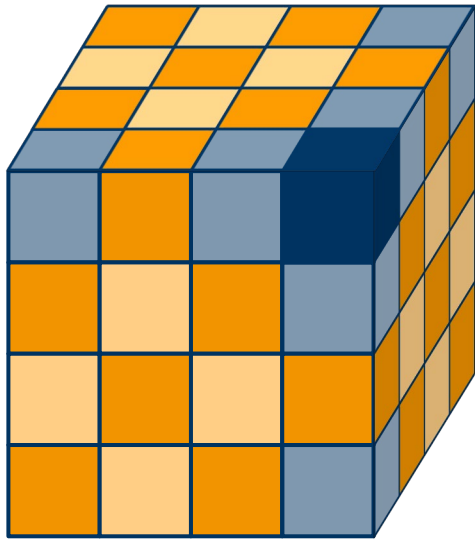
- ▶ cuts a grid into slices
 - ▶ a new communicator is generated for each slice
 - ▶ enables slice-restricted collective communication
- ▶ example on the right
 - ▶ slicing a 2D topology horizontally
 - ▶ 4 new communicators A, B, C, and D with 4 ranks each
 - ▶ `MPI_Bcast(..., A)` only affects ranks of A



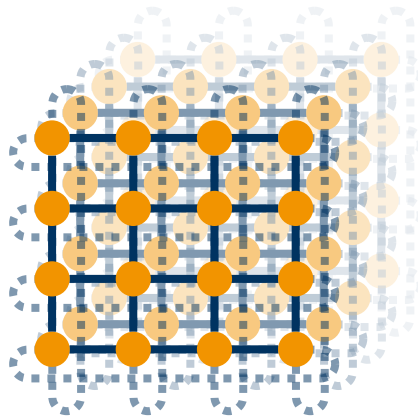
Convenience functions

- ▶ `MPI_Cart_coords(...)`
 - ▶ compute coordinates from a given rank (17 → [4, 1])
- ▶ `MPI_Cart_rank(...)`
 - ▶ compute rank from given coordinates ([4,1] → 17)
- ▶ `MPI_Cart_sub(...)`
 - ▶ partition grid into lower-dimensional sub-grids (e.g. 2D square from 3D cube)
- ▶ `MPI_Cartdim_get(...)/MPI_Cart_get(...)`
 - ▶ get topology information for a given communicator
- ▶ `MPI_Neighbor_allgather(...)/MPI_Neighbor_alltoall(...)`
 - ▶ sparse collective communication, exchanges data between neighbors if they are neighbors in the topology


Usefulness of Cartesian topologies



Use case: 3D heat stencil with periodic boundary conditions



Implementation: 3D Cartesian topology with periodicity in all three dimensions



Mapped to hardware that uses (at least a 3D) torus interconnect
(Image: Fugaku, Tofu D topology)

Tales from the Proseminar: Daniel's weird `int` problem

- ▶ sequential 2D heat stencil
 - ▶ 100x100 problem size
 - ▶ -O0 (but also with -O2) on LCC2
 - ▶ gcc 4.8.5 (but also 9.2 and MSVC 2019)
 - ▶ switching from `long` `long` loop iterators to `int` **halves** execution time

```
$ /usr/bin/time -f %E ./stencil_long 100 100
0:04.60
$ /usr/bin/time -f %E ./stencil_int 100 100
0:02.31
```

- ▶ What the heck is going on?

Tales from the Proseminar: Daniel's weird `int` problem

- ▶ profiled with gprof to find “hot spots”, left is `int`, right is `long long`

%	cumulative	self	
time	seconds	seconds	name
32.92	0.76	0.76	int.c:93
20.06	1.22	0.46	int.c:78
14.83	1.56	0.34	int.c:79
10.47	1.80	0.24	int.c:87
9.59	2.02	0.22	int.c:86

%	cumulative	self	
time	seconds	seconds	name
36.84	1.69	1.69	long.c:79
31.61	3.15	1.45	long.c:78
16.35	3.90	0.75	long.c:93
5.34	4.15	0.25	long.c:86
3.27	4.30	0.15	long.c:87

Tales from the Proseminar: Daniel's weird `int` problem cont'd

```
74 // get temperature at current position
75 value_t tc = A[i];
76
77 // get temperatures of adjacent cells
78 value_t txl = (i % Nx != 0) ? A[i - 1] : tc;
79 value_t txr = (i % Nx != Nx - 1) ? A[i + 1] : tc;
    // ..... snip .....
90 if (Ny > 1)
91     B[i] = tc + 0.165 * (txl + txr + tyl + tyr + tzl + tzt + (-6 * tc));
92 else
93     B[i] = tc + 0.2 * (txl + txr + tzl + tzt + (-4 * tc));
94 // if ((int)B[i] < (int)A[i])
```


Tales from the Proseminar: Daniel's weird `int` problem

- ▶ far-fetched idea, but maybe branch (miss-)predictions? compared with `perf stat`:

```
2328.650634 task-clock:u (msec) # 0.996 CPUs
          0 context-switches:u # 0.000 K/sec
          0 cpu-migrations:u   # 0.000 K/sec
        186 page-faults:u      # 0.080 K/sec
5,739,935,525 cycles:u         # 2.465 GHz
10,671,532,880 instructions:u  # 1.86 IPC
1,196,623,336 branches:u      # 513.870 M/sec
    1,030,678 branch-misses:u # 0.09%

2.338387625 seconds time elapsed
```

```
4639.728721 task-clock:u (msec) # 0.998 CPUs
          0 context-switches:u # 0.000 K/sec
          0 cpu-migrations:u   # 0.000 K/sec
        186 page-faults:u      # 0.040 K/sec
11,444,184,736 cycles:u        # 2.467 GHz
10,972,976,005 instructions:u  # 0.96 IPC
1,196,977,569 branches:u      # 257.984 M/sec
    1,030,347 branch-misses:u # 0.09%

4.650327844 seconds time elapsed
```

Tales from the Proseminar: Daniel's weird `int` problem cont'd

```
value_t tx1 = ( i % Nx != 0 ) ? A[i-1] : tc;
```

```
mov    eax, DWORD PTR [rbp-20]
cdq
idiv   DWORD PTR [rbp-24]
mov    eax, edx
test   eax, eax
je     .L2
mov    eax, DWORD PTR [rbp-20]
cdqe
```

```
sal    rax, 3
lea    rdx, [rax 8]
mov    rax, QWORD PTR [rbp-32]
add    rax, rdx
movsd  xmm0, QWORD PTR [rax]
jmp    .L3
.L2:   ...
.L3:   ...
```

Tales from the Proseminar: Daniel's weird `int` problem cont'd

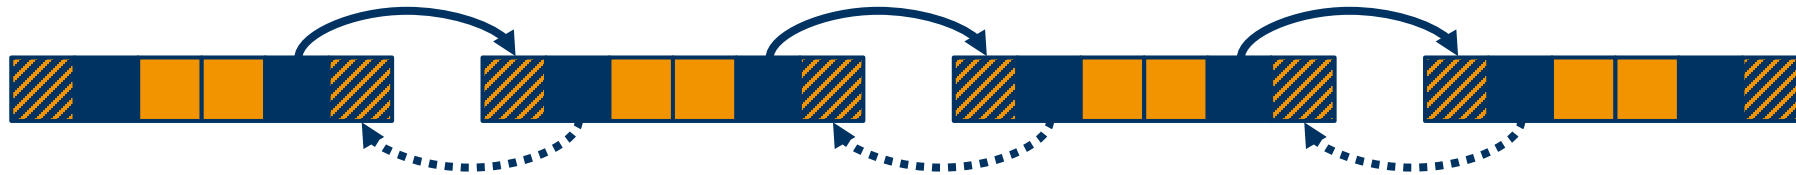
- ▶ Found instruction information on Agner Fog's blog for Intel Skylake architecture:
 - ▶ https://www.agner.org/optimize/instruction_tables.pdf
- ▶ if you want to study compiler output: <https://godbolt.org/>

inst.	operands	μops	latency
idiv	r32	10	26
idiv	r64	57	42-95

Tales from the Proseminar: Daniel's weird `int` problem cont'd

- ▶ root cause of the issues: single loop for multi-dimensional problem space
 - ▶ requires % operator to get boundaries (which are strided in linearized space)
 - ▶ replace with loop nest and comparison operators (>, <, ==, !=)
- ▶ potentially premature optimization and violates step 1 of “*Four Steps to Creating an Optimized Parallel Program*”

Tales from the proseminar: efficient ghost cell exchange

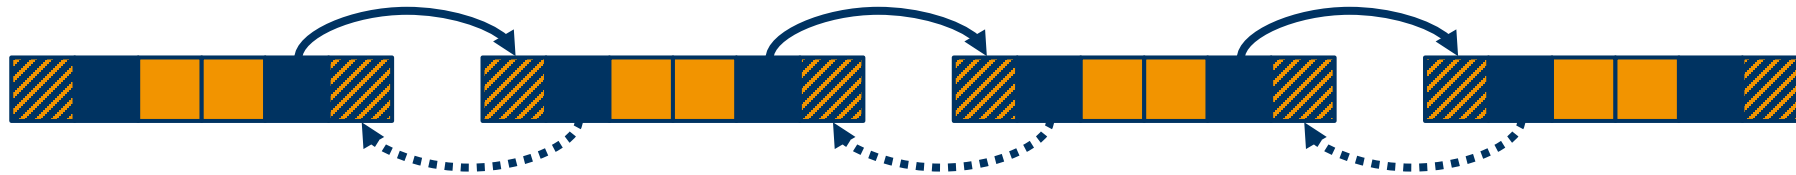


```
if(myRank > 0) {  
    // recv(left)  
}  
if(myRank < N-1) {  
    // send(right)  
}
```

time	rank 0	rank 1	rank 2	rank 3	...
0	send(right)	recv(left)	recv(left)	recv(left)	
1		send(right)	recv(left)	recv(left)	
2			send(right)	recv(left)	
3				send(right)	
...					

unnecessary serialization!

Tales from the proseminar: efficient ghost cell exchange cont'd



```
if(myRank % 2) {  
    // recv, send  
} else {  
    // send, recv  
}  
// or  
MPI_Sendrecv(...)
```

time	rank 0	rank 1	rank 2	rank 3	...
0	send(right)	recv(left)	send(right)	recv(left)	
1		send(right)	recv(left)	send(right)	

done!

Summary

- ▶ derived data types can be very handy
 - ▶ no need to copy data to basic, contiguous buffers
 - ▶ allows to easily transpose data
 - ▶ arbitrary nesting possible
- ▶ virtual topologies add semantic position information to ranks
 - ▶ makes rank positions easily identifiable
 - ▶ allows direct neighbor communication
 - ▶ enables limited-scope collectives
- ▶ Tales from the proseminar
 - ▶ Daniel's weird int problem
 - ▶ efficient ghost cell exchange

Image Sources

- ▶ Tofu D interconnect: https://link.springer.com/content/pdf/10.1007/978-3-319-07518-1_35.pdf